

3D Point Clouds

Lecture 4 – Clustering & Model Fitting

主讲人 黎嘉信

Aptiv 自动驾驶
新加坡国立大学 博士
清华大学 本科





1. Clustering - Spectral Clustering



2. Clustering – Mean-Shift and DBSCAN



3. Model Fitting - LSQ



4. Model Fitting – Hough Transform



5. Model Fitting - RANSAC



Lecture 3



K-Means

- Euclidean distance
- Hard assignment
- No modeling for a cluster
- Pre-defined cluster number k



GMM







- Euclidean distance
- Probability formulation – soft clustering
- Mean and variance estimation for each cluster
- Pre-defined cluster number k









Spectral Clustering

- Works with connectivity
- Heuristic to determine cluster number k






MiniBatchKMeans



GaussianMixture



SpectralClustering





Unnormalized Spectral Clustering

1. Build the graph to get adjacency matrix $W \in \mathbb{R}^{n \times n}$
2. Compute **unnormalized Laplacian L**
3. Compute the first (smallest) k eigenvectors **v_1, \dots, v_k of L**
4. Let $V \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors v_1, \dots, v_k as columns
5. For $i = 1, \dots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i -th row of V
6. Cluster the points $\{y_i \in \mathbb{R}^k\}$ with k-means algorithm into clusters C_1, \dots, C_k
7. The final output clusters are A_1, \dots, A_k where $A_i = \{j | y_j \in C_i\}$



Spectral Clustering

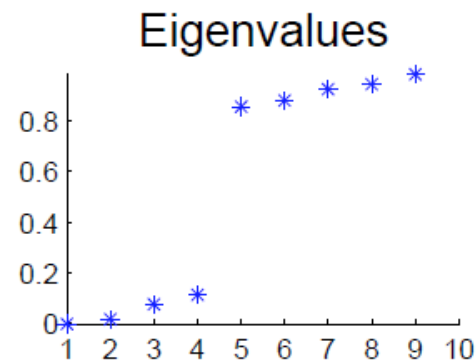
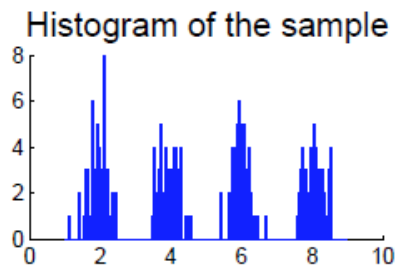


Selection of k can be done by eigenvalue analysis



Most stable clustering is given by the value of k that maximizes the eigen-gap

- Eigengap is the difference between consecutive eigenvalues
- $\Delta_k = |\lambda_k - \lambda_{k-1}|$





Normalized Spectral Clustering



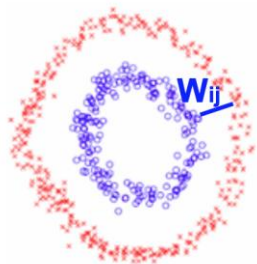
Unnormalized Spectral Clustering -> **approximated RatioCut**

$$\text{RatioCut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \bar{A}_i)}{|A_i|}$$

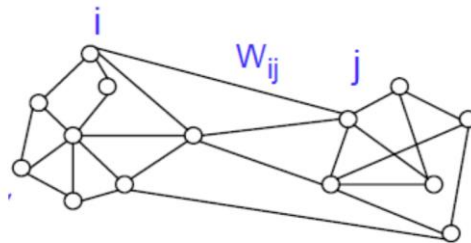


Normalized Spectral Clustering -> **approximated NormalizedCut**

$$\text{Ncut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \bar{A}_i)}{\text{vol}(A_i)}$$



Data clustering



$G = \{V, E\}$




Graph Min-cut




Min-cut: Partition graph $G = (V, E)$ into two sets A, B , such that the weights of edges connecting A, B is minimum

$$cut(A, B) = \sum_{i \in A, j \in B} w_{ij}$$



$G = \{V, E\}$



Min-cut for k partitions. \bar{A}_i is the complement of subset $A_i \subset V$

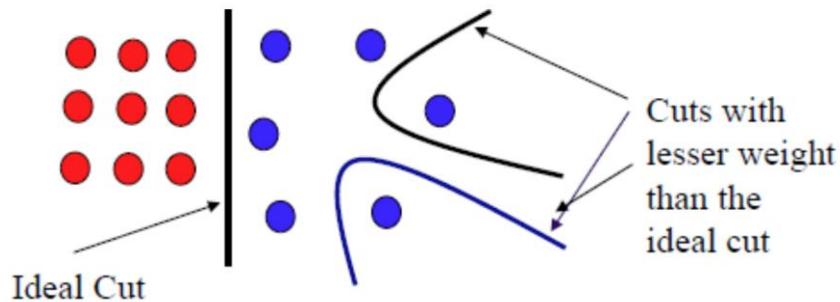
$$cut(A_1, \dots, A_k) = \sum_{i=1}^k cut(A_i, \bar{A}_i)$$



Graph Min-cut



Naïve min-cut suffers from degenerate results



Solution:

- Add constraints that the partition A_i can not be too small



How to evaluate the size of a partition A_i ?



Graph Min-cut – Size of A

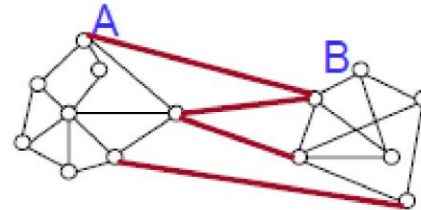
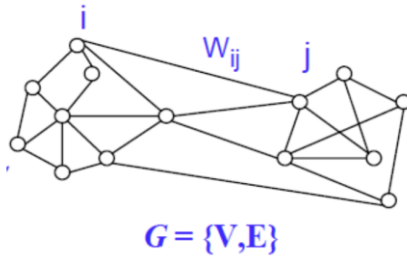
1. Number of vertices in A

$|A| :=$ the number of vertices in A

2. Sum of weights in A

$$\text{vol}(A) := \sum_{i \in A} d_i$$

$$d_i = \sum_{j=1}^n w_{ij} \quad \text{Weight sum for node } i$$





Spectral Clustering



Unnormalized Spectral Clustering -> **approximated RatioCut**

- Constrains the cluster have similar size
- $\text{Size}(A) = |A|$



Normalized Spectral Clustering -> **approximated NormalizedCut**

- Constrains the cluster have similar size
- $\text{Size}(A) = \text{vol}(A)$
- **$\text{Vol}(A)$ is large means nodes are closely connected inside A \rightarrow A's element is similar.**



Spectral Clustering – Laplacian Matrix



Degree matrix D

- A diagonal matrix with degrees d_1, \dots, d_n on the diagonal
- $d_i = \sum_{j=1}^n w_{ij}$ is the row sum of adjacency matrix $W \rightarrow$ “how many edges are connected to node i ”



Unnormalized graph Laplacian matrix $L = D - W$



Normalized graph Laplacian matrix

- $L_{sym} = D^{-1/2} L D^{-1/2} = I - D^{-1/2} W D^{-1/2}$
- $L_{rw} = D^{-1} L = I - D^{-1} W$

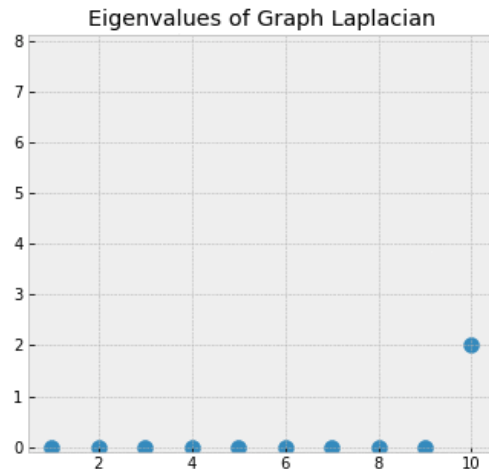
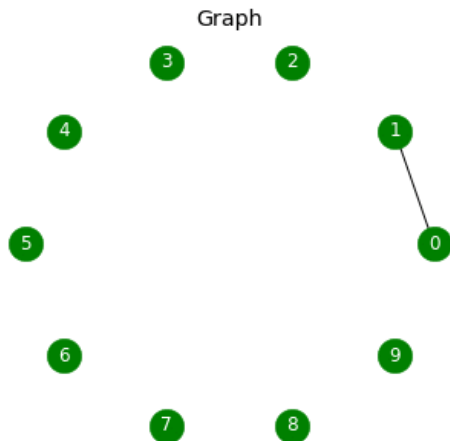


Spectral Clustering



Why do we need Laplacian matrix L ?

- Its eigenvalues / eigenvectors has some good properties.
 - One 0 eigenvalue \rightarrow one connected component
 - Corresponding eigenvectors \rightarrow which node belongs to that connected components







Spectral Clustering

$$W = \begin{bmatrix} 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$



Why do we need Laplacian matrix L ?

- Its eigenvalues / eigenvectors has some good properties.
 - One 0 eigenvalue \rightarrow one connected component
 - Corresponding eigenvectors \rightarrow which node belongs to that connected components



1	0
1	0
1	0
0	1
0	1
0	1
0	1
0	1
1	0
1	0

=

1	0
1	0
1	0
1	1
1	1
1	1
1	1
1	1
1	0
1	0



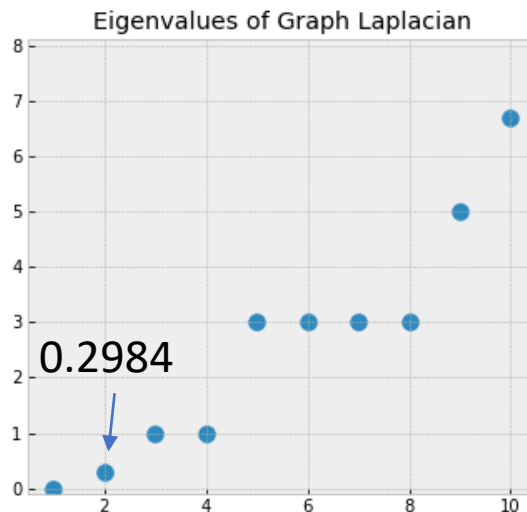
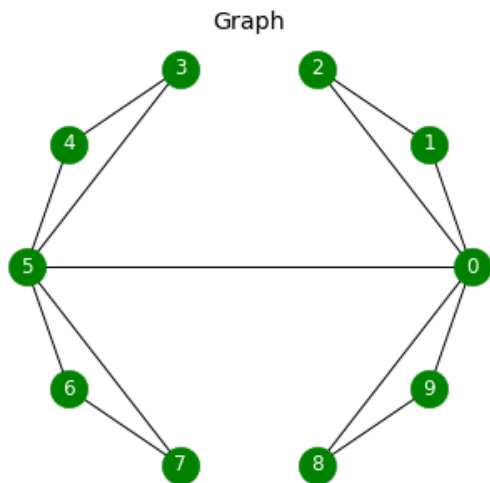
Spectral Clustering

0	1	1	0	0	1	0	0	1	1
1	0	1	0	0	0	0	0	0	0
1	1	0	0	0	0	0	0	0	0
0	0	0	0	1	1	0	0	0	0
0	0	0	1	0	1	0	0	0	0
1	0	0	1	1	0	1	1	0	0
0	0	0	0	0	1	0	1	0	0
0	0	0	0	0	1	1	0	0	0
1	0	0	0	0	0	0	0	0	1
1	0	0	0	0	0	0	0	1	0



Why do we need Laplacian matrix L ?

- Its eigenvalues / eigenvectors has some good properties.
 - One 0 eigenvalue \rightarrow one connected component
 - Corresponding eigenvectors \rightarrow which node belongs to that connected components



1	0
1	0
1	0
0	1
0	1
0	1
0	1
0	1
1	0
1	0



1	-0.23
1	-0.33
1	-0.33
1	0.33
1	0.33
1	0.23
1	0.33
1	0.33
1	-0.33
1	-0.33



Spectral Clustering – Graph Cut View



Unnormalized graph Laplacian matrix $L = D - W$

Proposition 1 (Properties of L) *The matrix L satisfies the following properties:*

1. *For every vector $f \in \mathbb{R}^n$ we have*

$$f^T L f = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{ij} (f_i - f_j)^2$$

2. *L is symmetric and positive semi-definite.*

3. *The smallest eigenvalue of L is 0, the corresponding eigenvector is the constant one vector $\mathbb{1}$.*

4. *L has n non-negative, real-valued eigenvalues $0 = \lambda_1 \leq \lambda_2 \leq \dots \leq \lambda_n$.*



Proposition 1.1

1. For every vector $f \in \mathbb{R}^n$ we have $f^T L f = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{ij} (f_i - f_j)^2$

$$f^T L f = f^T D f - f^T W f$$



Proposition 1.2 & 1.3 & 1.4

2. *L is symmetric and positive semi-definite.*

- $L = D - W$, D, W are symmetric
- $f^T L f \geq 0$, $\forall f \in \mathbb{R}^n$

3. *The smallest eigenvalue of L is 0, the corresponding eigenvector is the constant one vector $\mathbb{1}$.*

$$L f = (D - W) f = [\cdots, d_i f_i - \sum_{j=1}^n w_{ij} f_j, \cdots]^T = 0 \cdot f, \quad f = \mathbb{1}$$

$$d_i = \sum_{j=1}^n w_{ij}$$



4. *L has n non-negative, real-valued eigenvalues $0 = \lambda_1 \leq \lambda_2 \leq \cdots \leq \lambda_n$.*

Direct result of 2 & 3



Spectral Clustering – Graph Cut View

Proposition 2 (Number of connected components) *Let G be an undirected graph with non-negative weights. Then the multiplicity k of the eigenvalue 0 of L equals the number of connected components A_1, \dots, A_k in the graph. The eigenspace of eigenvalue 0 is spanned by the indicator vectors $\mathbb{1}_{A_1}, \dots, \mathbb{1}_{A_k}$ of those components.*



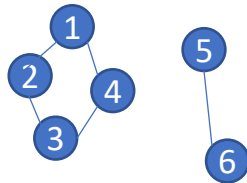


Spectral Clustering

- Wait, the eigenvector of the following graph is **NOT constant one**?

3. *The smallest eigenvalue of L is 0, the corresponding eigenvector is the constant one vector $\mathbb{1}$.*

- $Lx = \lambda x$
- $Lx_1 + Lx_2 = \lambda_1 x_1 + \lambda_2 x_2$
- $\lambda_1 = \lambda_2 = 0$
- $L(x_1 + x_2) = \lambda(x_1 + x_2)$



x_1	x_2
$\begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 0 \\ 0 \end{bmatrix}$	$\begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 1 \\ 1 \end{bmatrix}$

$$x_1 + x_2$$





Spectral Clustering – Graph Cut View

Proposition 2 (Number of connected components) *Let G be an undirected graph with non-negative weights. Then the multiplicity k of the eigenvalue 0 of L equals the number of connected components A_1, \dots, A_k in the graph. The eigenspace of eigenvalue 0 is spanned by the indicator vectors $\mathbb{1}_{A_1}, \dots, \mathbb{1}_{A_k}$ of those components.*

- Consider $k = 1$, i.e., G is a connected graph
- Assume f is eigenvector associated with eigenvalue 0

$$f^T L f = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{ij} (f_i - f_j)^2 = f^T \cdot 0 = 0$$

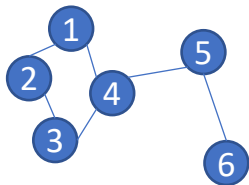
- $\forall i, j, f_i = f_j$ is the only choice to satisfy the above equation.



Spectral Clustering – Graph Cut View

$$f^T L f = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n w_{ij} (f_i - f_j)^2 = f^T \cdot 0 = 0$$

Graph



Connectivity / Similarity matrix

$$\begin{bmatrix} 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix}$$

$$f^T L f = (f_1 - f_2)^2 + (f_2 - f_3)^2 + (f_3 - f_4)^2 + (f_4 - f_5)^2 + (f_5 - f_6)^2 = 0$$

\Rightarrow

$$f_1 = f_2 = f_3 = f_4 = f_5 = f_6$$

That is, the smallest eigenvector is constant vector



Spectral Clustering – Graph Cut View

Proposition 2 (Number of connected components) *Let G be an undirected graph with non-negative weights. Then the multiplicity k of the eigenvalue 0 of L equals the number of connected components A_1, \dots, A_k in the graph. The eigenspace of eigenvalue 0 is spanned by the indicator vectors $\mathbb{1}_{A_1}, \dots, \mathbb{1}_{A_k}$ of those components.*

- Consider $k \geq 2$, reorganize L into **block diagonal matrix** on the left
- The spectrum of L is the union of spectrum of L_i . So the 0-eigenvalue eigenvectors of L is shown on the right

$$L = \begin{pmatrix} L_1 & & & \\ & L_2 & & \\ & & \ddots & \\ & & & L_k \end{pmatrix}$$

$$\begin{bmatrix} \mathbb{1}_1 & 0 & \cdots & 0 \\ 0 & \mathbb{1}_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & 0 \\ 0 & 0 & \cdots & \mathbb{1}_k \end{bmatrix}$$



Spectral Clustering – Graph Cut Definition

For two disjoint subsets, $A, B \subset V$

For k disjoint subsets, A_1, \dots, A_k

$$\text{cut}(A, B) = \sum_{i \in A, j \in B} w_{ij}.$$

$$\text{cut}(A_1, \dots, A_k) = \sum_{i=1}^k \text{cut}(A_i, \overline{A_i})$$

- Unnormalized Spectral Clustering -> **approximated RatioCut**

$$\text{RatioCut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \overline{A_i})}{|A_i|}$$

- Normalized Spectral Clustering -> **approximated NormalizedCut**

$$\text{Ncut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \overline{A_i})}{\text{vol}(A_i)}$$



Approximated RatioCut for $k = 2$

- The problem is simplified into

$$\min_{A \subset V} \text{RatioCut}(A, \bar{A}) = \min_{A \subset V} \left(\frac{\text{cut}(A, \bar{A})}{|A|} + \frac{\text{cut}(\bar{A}, A)}{|\bar{A}|} \right)$$

- Given a subset $A \subset V$, construct a vector $f = [f_1, \dots, f_n]^T \in \mathbb{R}^n$

$$f_i = \begin{cases} \sqrt{|\bar{A}|/|A|} & \text{if } v_i \in A \\ -\sqrt{|A|/|\bar{A}|} & \text{if } v_i \in \bar{A}. \end{cases}$$

- f determines the results of graph cut, but how to solve f ?

$$\begin{cases} v_i \in A & \text{if } f_i \geq 0 \\ v_i \in \bar{A} & \text{if } f_i < 0. \end{cases}$$



Approximated RatioCut for $k = 2$



Apply Proposition 1.1

$$f^T L f = \frac{1}{2} \sum_{i,j=1}^n w_{ij} (f_i - f_j)^2$$

$$= \frac{1}{2} \sum_{i \in A, j \in \bar{A}} w_{ij} \left(\sqrt{\frac{|\bar{A}|}{|A|}} + \sqrt{\frac{|A|}{|\bar{A}|}} \right)^2 + \frac{1}{2} \sum_{i \in \bar{A}, j \in A} w_{ij} \left(-\sqrt{\frac{|\bar{A}|}{|A|}} - \sqrt{\frac{|A|}{|\bar{A}|}} \right)^2$$

$$= \text{cut}(A, \bar{A}) \left(\frac{|\bar{A}|}{|A|} + \frac{|A|}{|\bar{A}|} + 2 \right)$$

$$= \text{cut}(A, \bar{A}) \left(\frac{|A| + |\bar{A}|}{|A|} + \frac{|A| + |\bar{A}|}{|\bar{A}|} \right)$$

$$= |V| \cdot \text{RatioCut}(A, \bar{A}).$$

$$f_i = \begin{cases} \sqrt{|\bar{A}|/|A|} & \text{if } v_i \in A \\ -\sqrt{|A|/|\bar{A}|} & \text{if } v_i \in \bar{A}. \end{cases}$$

$$\text{cut}(A, B) = \sum_{i \in A, j \in B} w_{ij}$$

$$\text{RatioCut}(A, \bar{A}) = \frac{\text{cut}(A, \bar{A})}{|A|} + \frac{\text{cut}(\bar{A}, A)}{|\bar{A}|}$$



Approximated RatioCut for $k = 2$



In addition, f is orthogonal to constant vector $\mathbb{1}$

$$\sum_{i=1}^n f_i = \sum_{i \in A} \sqrt{\frac{|\bar{A}|}{|A|}} - \sum_{i \in \bar{A}} \sqrt{\frac{|A|}{|\bar{A}|}} = |A| \sqrt{\frac{|\bar{A}|}{|A|}} - |\bar{A}| \sqrt{\frac{|A|}{|\bar{A}|}} = 0.$$



$$\|f\| = \sqrt{n}$$

$$\|f\|^2 = \sum_{i=1}^n f_i^2 = |A| \frac{|\bar{A}|}{|A|} + |\bar{A}| \frac{|A|}{|\bar{A}|} = |\bar{A}| + |A| = n.$$



Approximated RatioCut for $k = 2$



Now the problem is converted to:

$$\min_{ACV} f^T L f, \text{ s.t., } f \perp \mathbb{1}, \|f\| = \sqrt{n},$$

$$f_i = \begin{cases} \sqrt{|\bar{A}|/|A|} & \text{if } v_i \in A \\ -\sqrt{|A|/|\bar{A}|} & \text{if } v_i \in \bar{A}. \end{cases}$$



Approximation by dropping the last condition:

$$\min_{ACV} f^T L f, \text{ s.t., } f \perp \mathbb{1}, \|f\| = \sqrt{n}$$



Approximated RatioCut for $k = 2$



We are solving $\min_{A \subset V} f^T L f, \text{ s.t. }, f \perp \mathbb{1}, \|f\| = \sqrt{n}$



L is symmetric and Positive Semi-Definite



Recall the Rayleigh Quotients

Given a symmetric matrix $A \in S^n$,

$$\lambda_{\min}(A) \leq \frac{x^T A x}{x^T x} \leq \lambda_{\max}(A), \forall x \neq 0$$

$$\lambda_{\max}(A) = \max_{x: \|x\|_2=1} x^T A x$$

$$\lambda_{\min}(A) = \min_{x: \|x\|_2=1} x^T A x$$

The maximum and minimum are attained for $x = u_1$ and for $x = u_n$, respectively, where u_1 and u_n are the largest and smallest eigenvector of A , respectively.



Approximated RatioCut for $k = 2$





f is the **second** eigenvector of L



How can we get clusters?

$$\begin{cases} v_i \in A & \text{if } f_i \geq 0 \\ v_i \in \bar{A} & \text{if } f_i < 0. \end{cases}$$


$$f_i = \begin{cases} \sqrt{|\bar{A}|/|A|} & \text{if } v_i \in A \\ -\sqrt{|A|/|\bar{A}|} & \text{if } v_i \in \bar{A}. \end{cases}$$




Run k-means on f .

- In practice, run k-means on the first two eigenvectors
- First eigenvector is graph connectivity, which helps as well


input affinity matrix



affinity matrix reordered according to solution vector



the partition according to the solution vector





Approximated RatioCut for $k \geq 2$

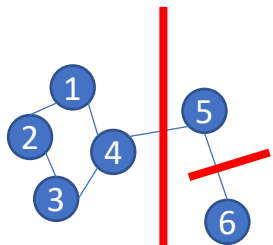


Construction the indication matrix $H \in \mathbb{R}^{n \times k}$, where,

$$h_{i,j} = \begin{cases} 1/\sqrt{|A_i|} & \text{if } i \in A_j \\ 0 & \text{otherwise.} \end{cases}$$



A vertex belongs one cluster only $\rightarrow H$ has orthonormal columns $H^T H = I$



$$H = \begin{bmatrix} \frac{1}{\sqrt{4}} & 0 & 0 \\ \frac{1}{\sqrt{4}} & 0 & 0 \\ \frac{1}{\sqrt{4}} & 0 & 0 \\ \frac{1}{\sqrt{4}} & 0 & 0 \\ 0 & \frac{1}{\sqrt{1}} & 0 \\ 0 & 0 & \frac{1}{\sqrt{1}} \end{bmatrix}$$



Approximated RatioCut for $k \geq 2$



Construction the indication matrix $H \in \mathbb{R}^{n \times k}$, where,

$$h_{i,j} = \begin{cases} 1/\sqrt{|A_i|} & \text{if } i \in A_j \\ 0 & \text{otherwise.} \end{cases}$$



A vertex belongs one cluster only $\rightarrow H$ has orthonormal columns $H^T H = I$



Denote the row vector as $h_i \in \mathbb{R}^k, i = 1, \dots, n$



Following similar calculations as $k = 2$, we have,

$$h_i^T L h_i = \frac{\text{cut}(|A_i|, |\bar{A}_i|)}{|A_i|}, \quad h_i^T L h_i = (H^T L H)_{ii}$$



Approximated RatioCut for $k \geq 2$



Recall the definition of RatioCut, we have

$$\begin{aligned}\text{RatioCut}(A_1, \dots, A_k) &= \sum_{i=1}^k \frac{\text{cut}(A_i, \bar{A}_i)}{|A_i|} \\ &= \sum_{i=1}^k h_i^T L h_i = \sum_{i=1}^k (H^T L H)_{ii} = \text{Tr}(H^T L H)\end{aligned}$$



Now the problem is transformed into

$$\min_{A_1, \dots, A_k} \text{Tr}(H^T L H) \text{ s.t., } H^T H = I, \quad h_{i,j} = \begin{cases} 1/\sqrt{|A_i|} & \text{if } i \in A_j \\ 0 & \text{otherwise.} \end{cases}$$



Approximated RatioCut for $k \geq 2$



Approximation by dropping our H construction,

$$\min_{A_1, \dots, A_k} \text{Tr}(H^T L H) \text{ s.t., } H^T H = I$$



A more general form of Rayleigh Quotients, gives the solution:

- H contains the first k eigenvectors of L as columns



Apply k-means on the rows of H because of the approximation

~~$$h_{i,j} = \begin{cases} 1/\sqrt{|A_j|} & \text{if } i \in A_j \\ 0 & \text{otherwise.} \end{cases}$$~~



Normalized Spectral Clustering



The derivation is similar



The goal is Normalized Cut on the graph



Unnormalized Spectral Clustering -> **approximated RatioCut**

$$\text{RatioCut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \bar{A}_i)}{|A_i|}$$



Normalized Spectral Clustering -> **approximated NormalizedCut**

$$\text{Ncut}(A_1, \dots, A_k) = \sum_{i=1}^k \frac{\text{cut}(A_i, \bar{A}_i)}{\text{vol}(A_i)}$$



Intuition of Spectral Cluster

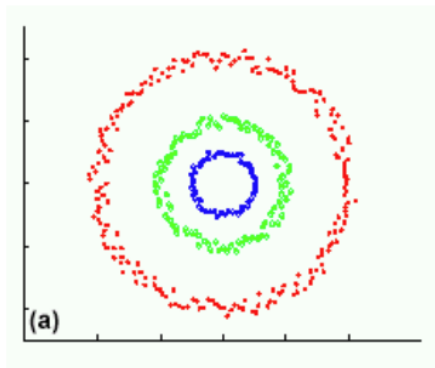
1. RatioCut or NormalizedCut

- Cut the graph into disjoint subsets with minimum cutting weights
- The size of the each subset should not be too small

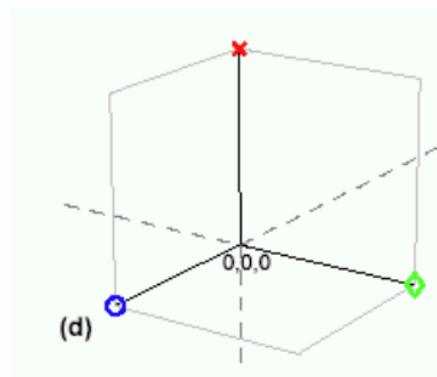
2. Dimension Reduction

- Project the n-dimension data into lower dimension space (spectral/eigenvector domain)

Original data



Projected data





Spectral Clustering - Summary



Complexity: $O(n^3)$

- This is the complexity of eigen decomposition
- K-means complexity is $O(n^2)$



Advantage

- No assumption on cluster shape
- Works with similarity, including Euclidean, connectivity
- Works with any dimensional data
- Able to estimate the number of clusters



Disadvantage

- Computational expensive
 - Can be alleviated using sparse similarity matrix and sparse eigen solver

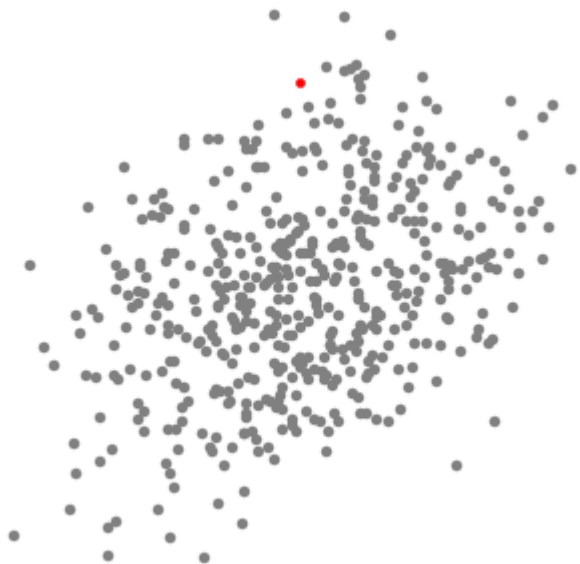


Mean Shift



Sliding windows hill climbing

- “Hill” is density

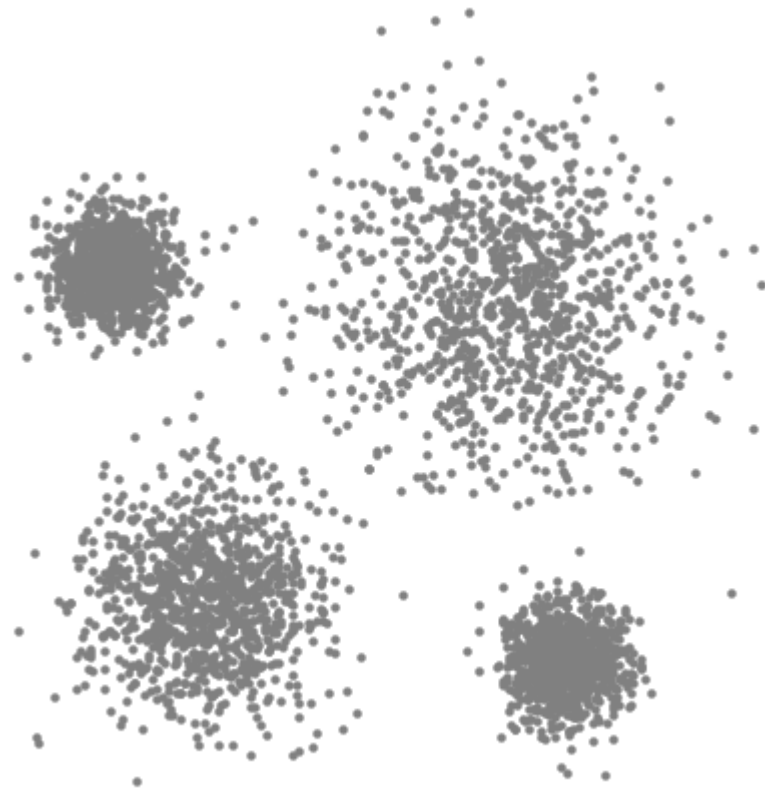




Mean Shift

1. Randomly select a circle with radius r
2. Move the circle to the center of the points inside
3. Repeat step-2 until it doesn't move
4. Repeat step-1,2,3. Remove overlapping circles
 - If circles overlap, select the one with most points
5. Determine clusters by finding the nearest circle center (similar to k-means)

Parameters: radius r





Mean Shift – Summary



Complexity $O(T \cdot n \cdot \log(n))$

- T is number of centers
- $n \cdot \log(n)$ is the complexity of radius based neighbor search, given 2D/3D data with kd-tree/octree



Advantage

- Automatically determines cluster numbers
- Single parameter
- Robust to outliers




Disadvantage

- Hill climbing easily falls in local minima
- Depends on initialization
- Assumes clusters are in ellipse shape
- Mainly works in Euclidean space
- Doesn't scale with high dimensional data



Density-Based Spatial Clustering of Applications with Noise (DBSCAN)





DBSCAN


Preparation: all points labeled as unvisited

Parameters: *distance r* , *min_samples*

1. Randomly select a unvisited point p , find its neighborhood within r
2. Number of points within $r \geq \text{min_samples}$?
 - Yes. p is a *core point*, Create a cluster C , go to step 3, mark p as *visited*.
 - No. Mark p as *noise* and *visited*.
3. Go through points within its r -neighborhood, label it as C
 - If it is a *core point*, set it as the “new p ”, repeat step-3
4. Remove cluster C from the database, go to step-1
5. Terminate when all points are visited.



DBSCAN



Red: Core points. point number within circle ≥ 4

Yellow: Border points. Still part of the cluster because it is within r of a core point, but does not meet the `min_points` criteria

Blue: Noise point. Not assigned to a cluster.



DBSCAN – Summary



Complexity $O(n \cdot \log(n))$

- Radius NN search for each point



Advantage

- No assumption on cluster shape
- Automatically determines cluster numbers
- Robust to outliers



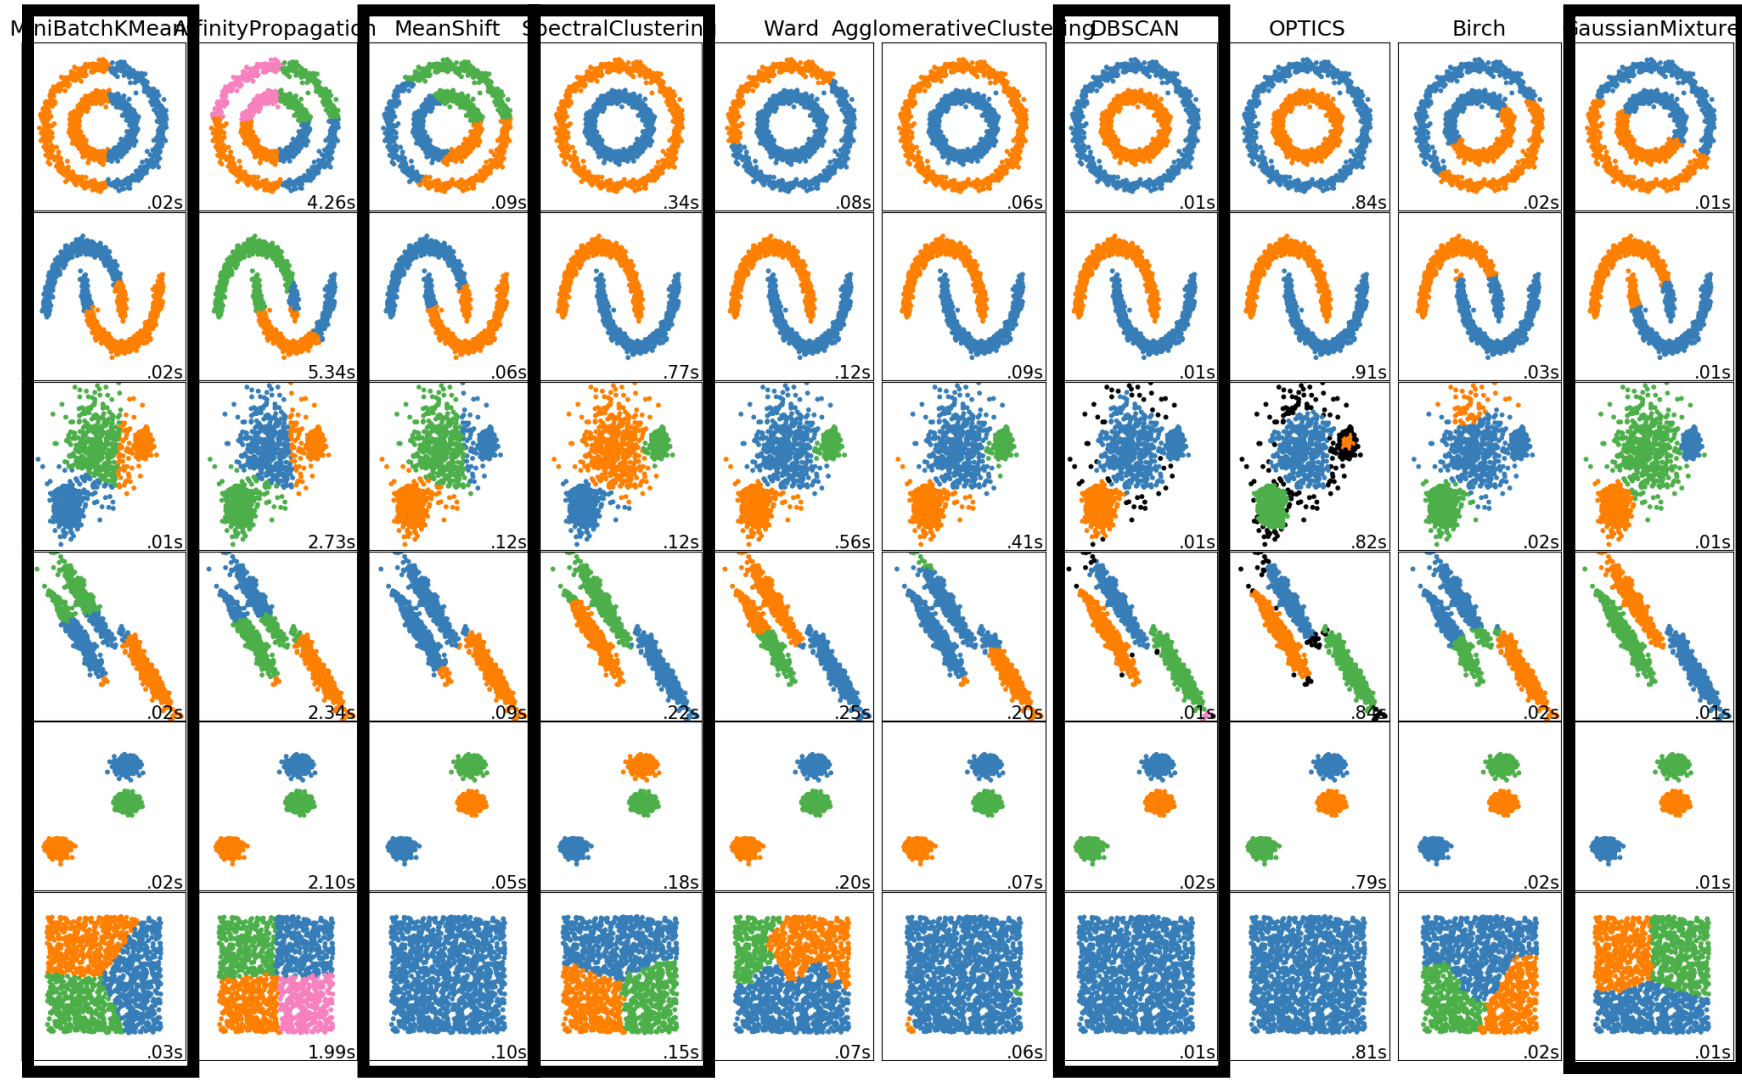
Disadvantage

- Doesn't work well with varying density
 - DBSCAN assume high density clusters are separated by some low density points
- Mainly works in Euclidean space
- Doesn't scale with high dimensional data



Clustering - Summary

	K-Means	GMM	Spectral	Mean Shift	DBSCAN
Metric	Euclidean	Euclidean	Similarity	Density /Euclidean	Density /Euclidean
# of clusters	Pre-defined	Pre-defined	Heuristic	Automatic	Automatic
Robustness to outlier	Bad	Medium	Good	Good	Good
High dimension data	Medium	Medium	Good	Bad	Bad
Complexity	$O(t \cdot k \cdot n \cdot d)$ t: iteration k: # of clusters n: # of data d: dimension	$O(t \cdot k \cdot n \cdot d)$ t: iteration k: # of clusters n: # of data d: dimension	$O(n^3)$ n: # of data	$O(Tn \log(n))$ n: # of data T: # of centers	$O(n \cdot \log(n))$ n: # of data





Model Fitting



Take line fitting as example



Approaches:

- Least Square
- Hough Transform
- Random Sample Consensus (RANSAC)



Model Fitting



If we know the inlier points

- Least Square



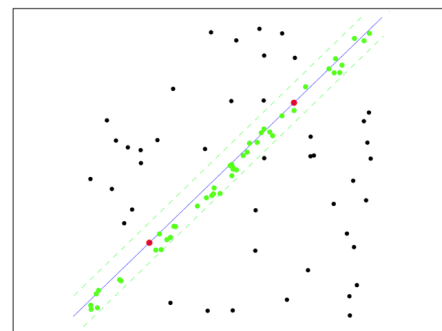
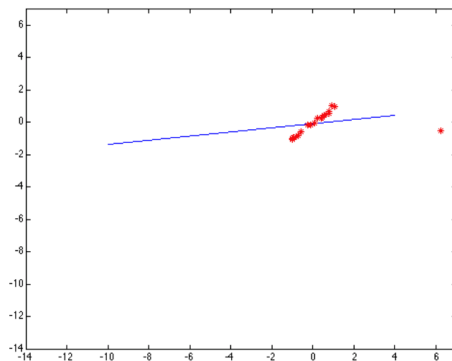
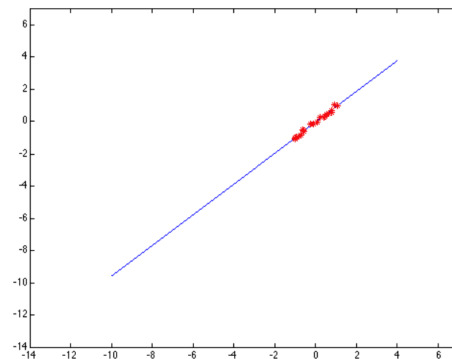
What if there is small amount of outliers?

- Robust Least Square, e.g., robust loss function
- Hough Transform
- RANSAC



What if there are lots of outliers / more than one models in data?

- Hough Transform
- RANSAC





Least Square Fitting



Given a set of points $\{p_1, \dots, p_n\}$, find a line that fits the point set best.

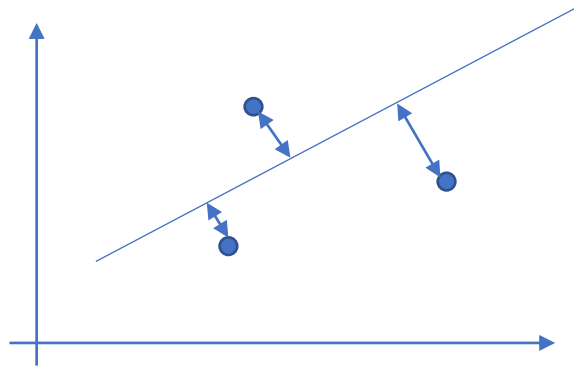


Consider 2-dimension, $p_i \in \mathbb{R}^2$, line model is $ax + by + c = 0$



Definition of “fit” – minimize the perpendicular distance

$$E = \sum_{i=1}^n (ax_i + by_i + c)^2$$





Least Square Fitting



Re-organize the problem into standard least-square optimization (sometimes called homogeneous equations):

$$\hat{\mathbf{x}} = [\hat{a}, \hat{b}, \hat{c}]^T = \min_{\mathbf{x}} \|\mathbf{A}\mathbf{x}\|_2^2, \text{ s.t. } \|\mathbf{x}\|_2 = 1$$

$$\mathbf{A} = \begin{bmatrix} x_1 & y_1 & 1 \\ \vdots & \vdots & \vdots \\ x_n & y_n & 1 \end{bmatrix}, \mathbf{x} = [a, b, c]^T$$



The solution is obvious: $[a, b, c]^T$ is the **eigenvector of the smallest eigenvalues of \mathbf{A}**

- Given \mathbf{A} is full column rank, i.e., $n \geq 3$



Least Square Fitting

Many model fitting problem can be formulated as least square (LSQ) optimization problem.

Linear LSQ problem $A\mathbf{x} = 0$

$$\hat{\mathbf{x}} = \min_{\mathbf{x}} \|A\mathbf{x}\|_2^2, \text{ s.t., } \|\mathbf{x}\|_2 = 1, A \in \mathbb{R}^{n \times m}, \mathbf{x} \in \mathbb{R}^m$$

- Solution given by eigenvector of the smallest eigenvalue of A

Linear LSQ problem $A\mathbf{x} = \mathbf{b}$:

$$\hat{\mathbf{x}} = \min_{\mathbf{x}} \|A\mathbf{x} - \mathbf{b}\|_2^2, A \in \mathbb{R}^{n \times m}, \mathbf{x} \in \mathbb{R}^m, \mathbf{b} \in \mathbb{R}^n$$

- In the case that $n \geq m$, the solution is given by $\hat{\mathbf{x}} = (A^T A)^{-1} A^T \mathbf{b}$

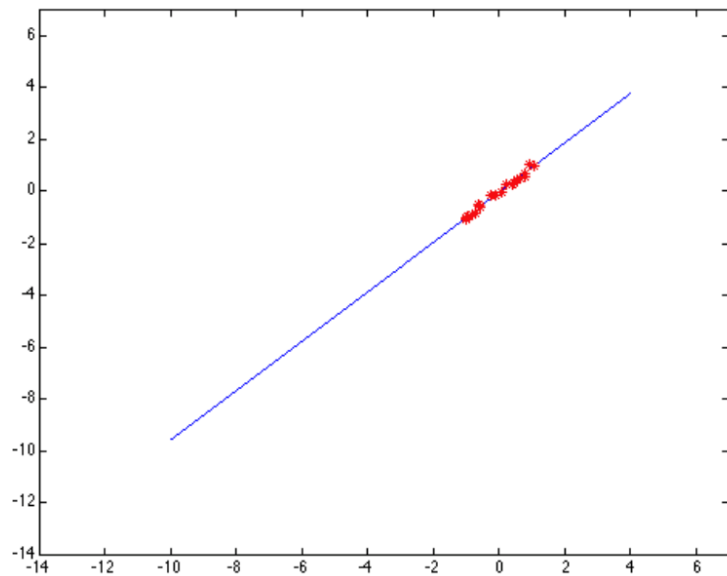
Linear LSQ problem $A\mathbf{x} = \mathbf{b}$, s.t., $C\mathbf{x} = 0$



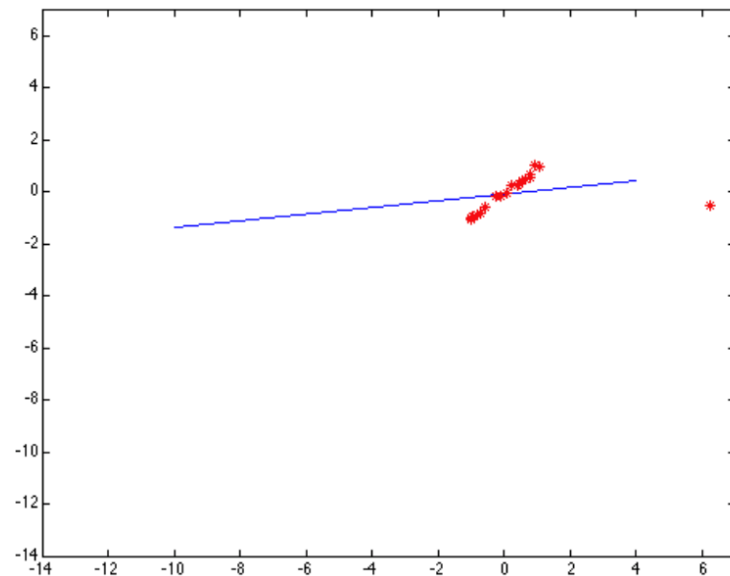
Limitations of LSQ



Sensitive to Outlier



LSQ with No Outlier



LSQ with One Outlier

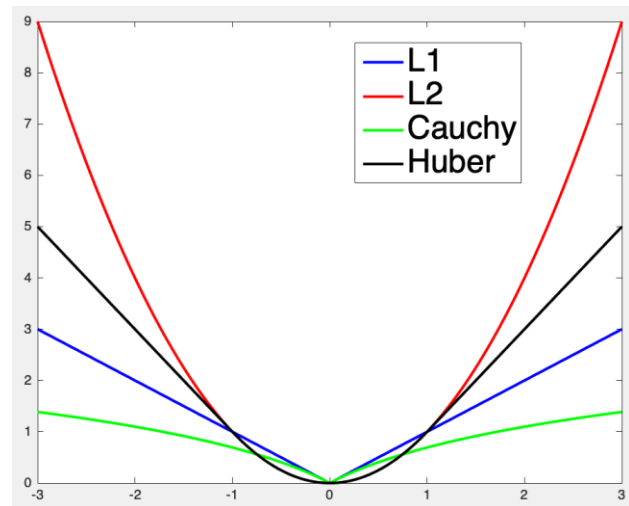


Loss functions



Typical loss functions

- L1. $\rho = |s|$
- L2. $\rho = s^2$
- Cauchy. $\rho = \log(1 + |s|)$
- Huber. $\rho = \begin{cases} s^2, & |s| < \delta \\ 2\delta(|s| - \frac{1}{2}\delta), & \text{otherwise} \end{cases}$
- etc.



Robust loss functions like Huber, Cauchy reduce the effect of outliers



However, the problem becomes *non-linear*!



Non-Linear LSQ



A general formulation of LSQ

$$\hat{\mathbf{x}} = \min_{\mathbf{x}} \|f(\mathbf{x})\|^2$$



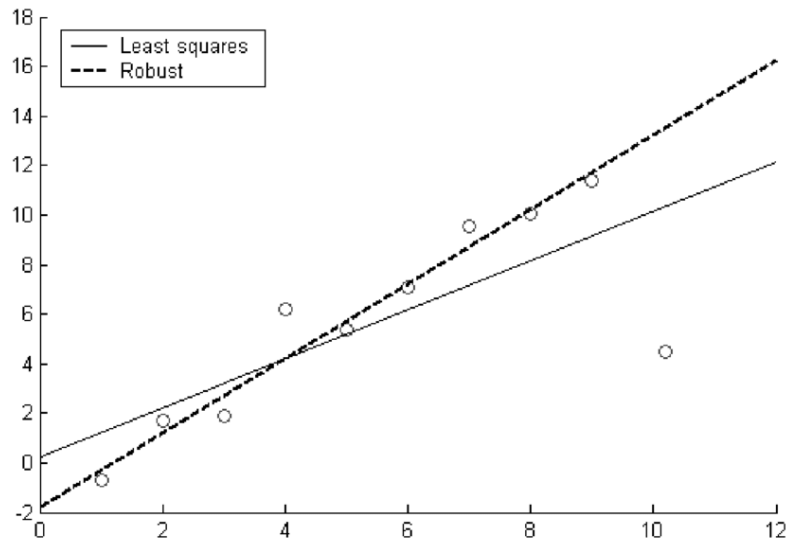
Function f is the non-linear function

- e.g., coupling the robust loss function with linear LSQ



Optimization methods

- Gradient descent
- Gauss-Newton
- Levenberg-Marquardt





Model Fitting



Least Square / Robust Least Square

- No outlier / few outlier
- Simple and fast



What if:

- Lots of outliers
- More than one models in the data, e.g., more than one line.



Hough Transform



Random Sample Consensus (RANSAC)



Hough Transform



Discretize parameter spaces into bins



For each data point, vote the bins that can generate this data point



Find the bins with most votes





Image space



Line parameter space



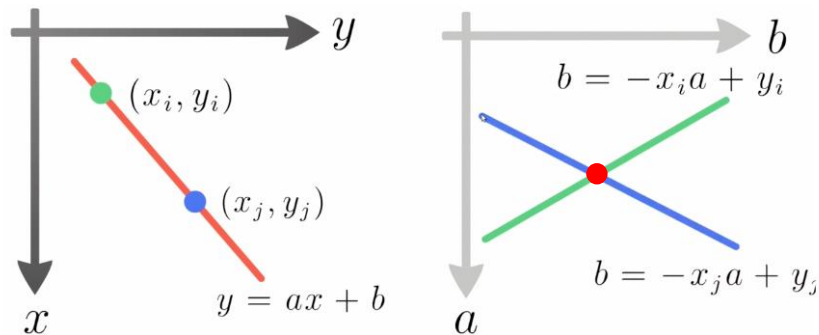
Hough Transform



A point in the Euclidean space \rightarrow A line in the parameter space



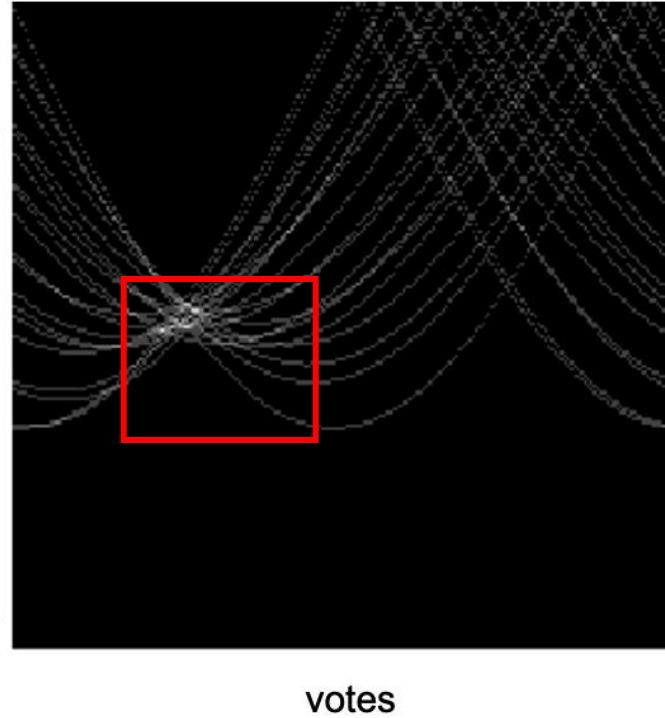

Select the bin with most votes





Hough Transform – Effect of Noise

Blurred peak
– in accurate parameter estimation





Hough Transform



Model parameterization. E.g., for a line

- $y = ax + b$ is non-uniform, can't represent vertical lines (a is infinity)
- $x\cos\theta + y\sin\theta = r$ is a better model with parameters $\{\theta, r\}$



Selection of resolution

- Tradeoff between speed and precision



Apply smoothing at the parameter space before searching for the highest vote

- E.g., Gaussian smooth
- Reduce the effect of noise



Hough Transform – Extension for Circles



The circle model with parameters $\{a, b, r\}$

$$(x - a)^2 + (y - b)^2 = r^2$$



Each point (x, y) fills a set of parameter bins for $\{a, b, r\}$ that fulfills the above function.

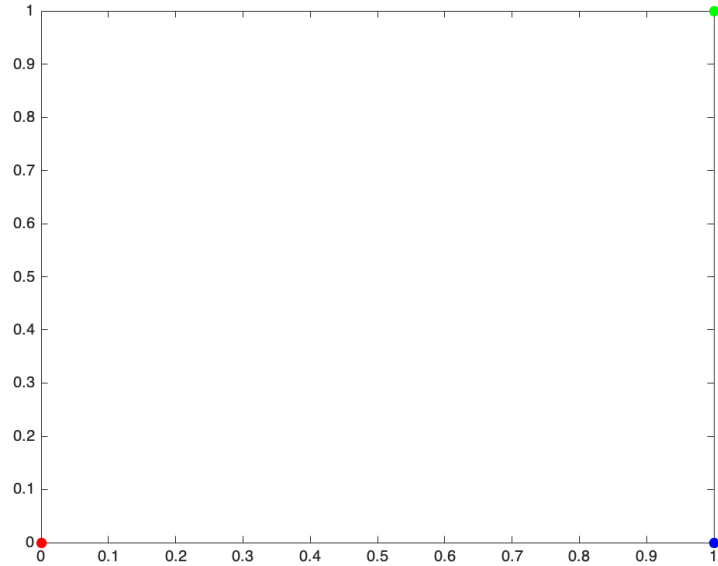


How to find those bins?

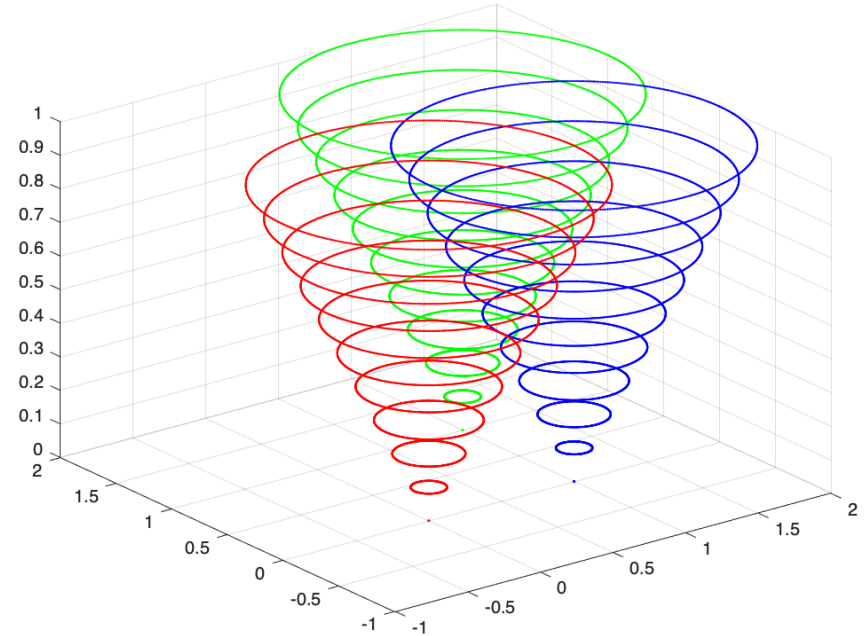
- Fix $r = r_i$
- Uniformly sample a set of $\{\theta_1, \dots, \theta_k\}$
- Each θ_j generates a set of $\{a, b\}$
 - $a = x - r_i \cos \theta_j$
 - $b = y - r_i \sin \theta_j$



Hough Transform – Extension for Circles



3 points in Euclidean space: $[0, 0]$, $[1, 1]$, $[1, 0]$



Each point corresponds to a "cone" in Hough space



Hough Transform – Summary



Advantage

- Robust to noise
- Robust to missing points of the shape
- Can be extended to lots of models



Disadvantage

- Doesn't scale well with complicated models
 - Usually works for models with less than 3 unknown parameters



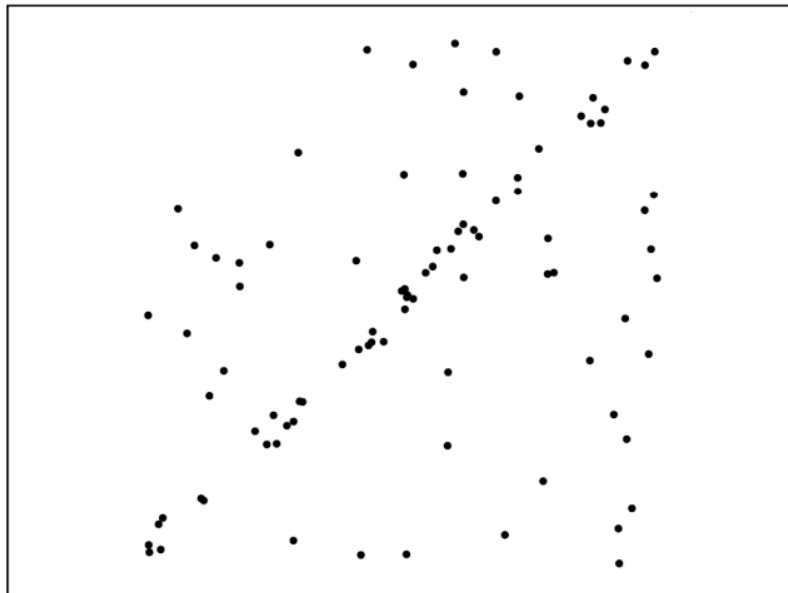
RANSAC



Simple and works well in practice



Works well with complicated models





RANSAC – Line Fitting

1. Randomly select a sample (minimal subset of points required to solve the model)

$$p_0 = (x_0, y_0), p_1 = (x_1, y_1)$$

2. Solve the line model

$$x = x_0 + at$$

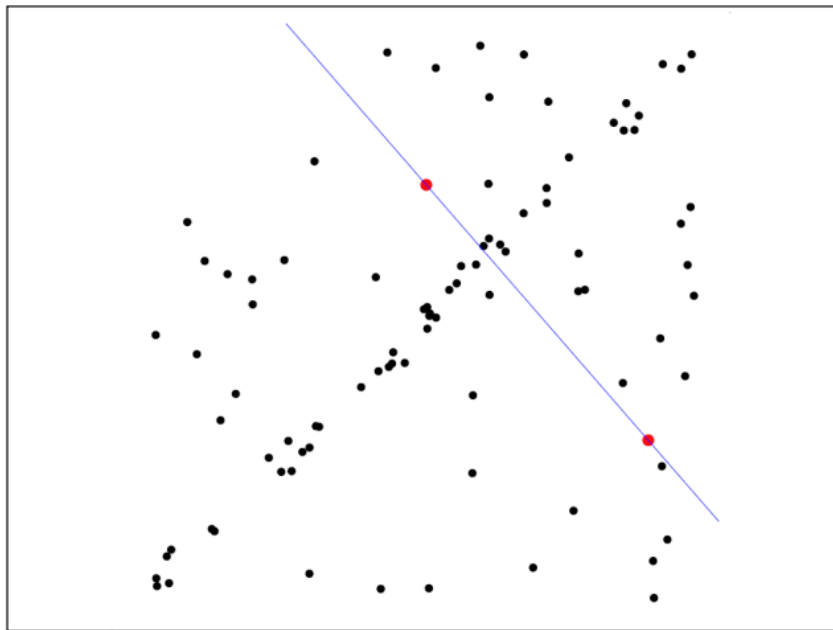
$$y = y_0 + bt$$

where,

$$p_0 = [x_0, y_0]^T, n = [a, b]^T$$

$$\Delta x = x_1 - x_0, \Delta y = y_1 - y_0$$

$$\Delta x = at, \Delta y = bt \rightarrow \frac{\Delta x}{\Delta y} = \frac{a}{b}$$






RANSAC – Line Fitting

1. Randomly select a minimal subset of points required to solve the model
2. Solve the model
3. Compute error function for each point

$$p_i = (x_i, y_i)$$

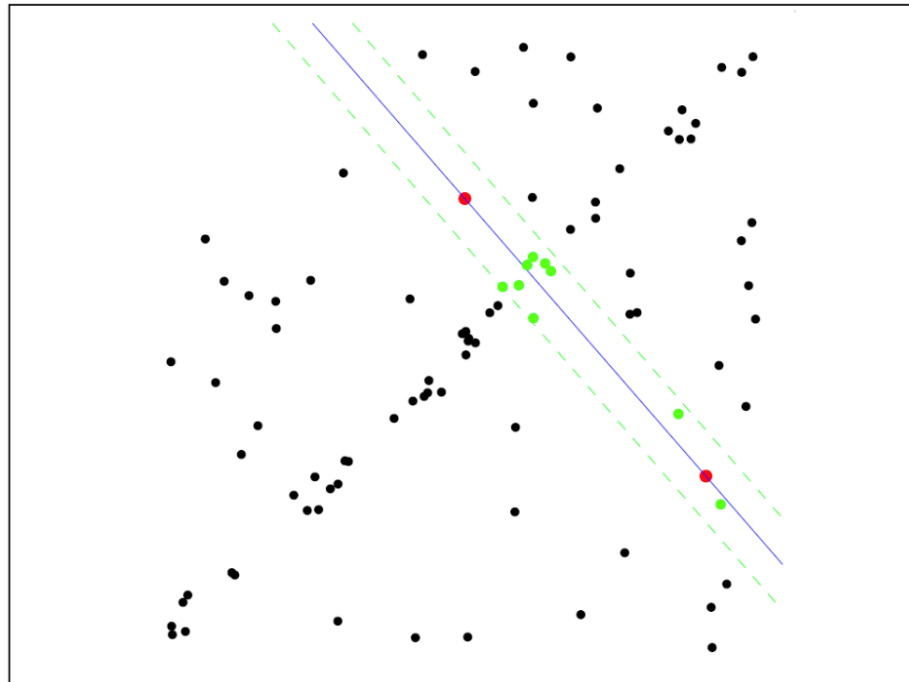
$$d_i = \frac{n^T(p_i - p_0)}{\|n\|_2}$$





RANSAC – Line Fitting

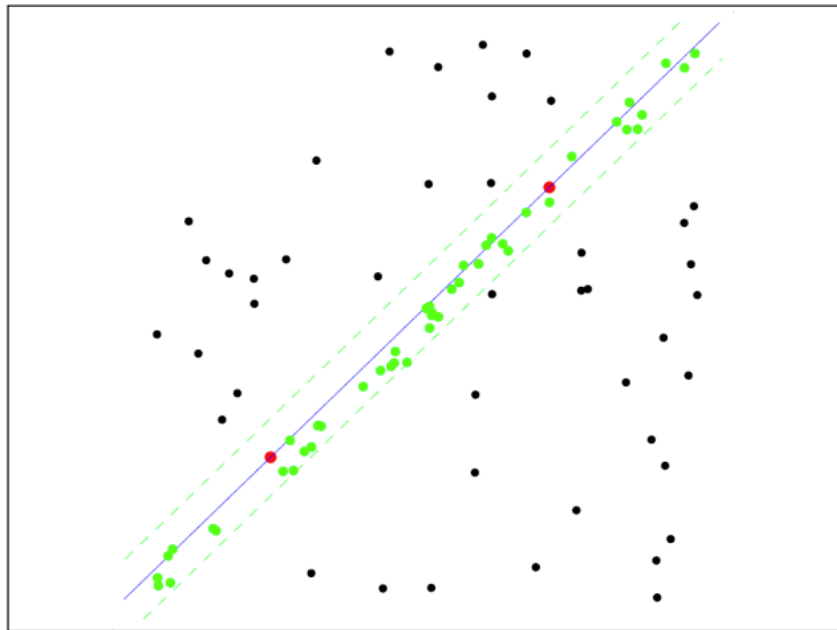
1. Randomly select a minimal subset of points required to solve the model
2. Solve the model
3. Compute error function for each point $p_i = (x_i, y_i)$
4. Count the points consistent with the model, $d_i < \tau$ (inlier)





RANSAC – Line Fitting

1. Randomly select a minimal subset of points required to solve the model
2. Solve the model
3. Compute error function for each point $p_i = (x_i, y_i)$
4. Count the points consistent with the model, $d_i < \tau$
5. Repeat step 1-4 for N iterations, choose the model with most inlier points





RANSAC – Line Fitting



Distance threshold τ

- Usually chosen empirically
- Chi-square distribution χ^2



Number of iterations N

- Choose N so that with probability p , as least one random sample is free from outliers, e.g., $p = 0.99$



RANSAC – Threshold τ by χ^2



Assume the error between data point and model is Gaussian distribution $d \sim \mathcal{N}(0, \sigma^2)$



χ^2 distribution – sum of squares of k independent standard normal distribution. Assume 95% confidence that the point is an inlier

- 1DoF χ_1^2 for 2D/3D line fitting or 3D plane fitting. The error is perpendicular distance
 - $\tau = \sqrt{3.84\sigma^2}$
- 2DoF χ_2^2 for 2D point distance, because the distance is $\Delta x^2 + \Delta y^2$
 - $\tau = \sqrt{5.99\sigma^2}$
- 3DoF χ_3^2 for 3D point distance, because the distance is $\Delta x^2 + \Delta y^2 + \Delta z^2$
 - $\tau = \sqrt{7.81\sigma^2}$



RANSAC – Threshold τ by χ^2

Degrees of freedom (df)	χ^2 value ^[19]											
1	0.004	0.02	0.06	0.15	0.46	1.07	1.64	2.71	3.84	6.63	10.83	
2	0.10	0.21	0.45	0.71	1.39	2.41	3.22	4.61	5.99	9.21	13.82	
3	0.35	0.58	1.01	1.42	2.37	3.66	4.64	6.25	7.81	11.34	16.27	
4	0.71	1.06	1.65	2.20	3.36	4.88	5.99	7.78	9.49	13.28	18.47	
5	1.14	1.61	2.34	3.00	4.35	6.06	7.29	9.24	11.07	15.09	20.52	
6	1.63	2.20	3.07	3.83	5.35	7.23	8.56	10.64	12.59	16.81	22.46	
7	2.17	2.83	3.82	4.67	6.35	8.38	9.80	12.02	14.07	18.48	24.32	
8	2.73	3.49	4.59	5.53	7.34	9.52	11.03	13.36	15.51	20.09	26.12	
9	3.32	4.17	5.38	6.39	8.34	10.66	12.24	14.68	16.92	21.67	27.88	
10	3.94	4.87	6.18	7.27	9.34	11.78	13.44	15.99	18.31	23.21	29.59	
P value (Probability)	0.95	0.90	0.80	0.70	0.50	0.30	0.20	0.10	0.05	0.01	0.001	

https://en.wikipedia.org/wiki/Chi-squared_distribution



RANSAC – Number of Iterations N



e : outlier ratio (probability that a point is an outlier)



s : number of points in a sample (e.g., in line fitting a sample contains 2 points)



N : sample number N (number of RANSAC iteration)



p : confidence we get at least a good sample that is free from outliers

$$\underbrace{(1 - (1 - e)^s)}^{\text{Probability of choosing } s \text{ inliers in a row}}^N = 1 - p \quad \underbrace{(1 - (1 - e)^s)}^{\text{Probability that one or more points are outliers}}^N = 1 - p \quad \underbrace{(1 - (1 - e)^s)}^{\text{Probability that } N \text{ samples are contaminated}}^N = 1 - p$$

Probability of choosing s
inliers in a row

Probability that one or more
points are outliers

Probability that N samples are
contaminated



RANSAC – Number of Iterations N



Iteration number N is given by

$$N = \frac{\log(1 - p)}{\log(1 - (1 - e)^s)}$$



Table for $p = 0.99$

proportion of outliers e							
s	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177



RANSAC - Practical Tricks




Don't need to perform N samples/iterations

- Terminate when the inlier ratio reach the expected inlier ratio

$$T = (1 - e) \cdot total_num_of_data_points$$



Run LSQ to refine the model after selecting the final model and inlier points





RANSAC - Summary



Advantages

- Simple and general
- Usually works well in practice, even with low inlier ratio like 10%



Disadvantages

- Need to determine the inlier threshold τ
- Need large number of samples when inlier ratio is low



Homework



Object detection pipeline for lidar

- Use KITTI 3D object detection dataset, select 3 point clouds, do the followings.
- Step 1. Remove the ground from the lidar points. Visualize ground as blue.
 - Any method you want – LSQ, Hough, RANSAC
- Step 2. Clustering over the remaining points. Visualize the clusters with random colors.
 - Any method you want
- Step 3. Classification over the clusters
 - Homework of Lecture 5
- Step 4. Report the detection precision-recall for three categories: vehicle, pedestrian, cyclist
 - Homework of Lecture 5